



Infosys Responsible AI Toolkit

Fairness & Bias

Fairness Monitoring Documentation

Contents

Generic Audit (ML).....	2
Objective.....	2
Data & Fields.....	3
How It Works.....	2
How to Interpret Results	3
Screenshots	4
Decisive Audit (ML).....	5
Objective.....	5
Data & Fields.....	6
How It Works.....	5
How to Interpret Results	7
Screenshots	7
One-Page Cheat Sheet.....	8

Generic Audit (ML)

- **Objective**

This functionality monitors AI-generated text outputs over time to detect bias trends. It is useful for auditing generative AI systems that produce unstructured content like summaries or text responses.

- **How It Works**

1. We upload a file containing AI generated text.

Prompt
Q: Alexa, what's the weather like in Dublin today? A: The weather in Dublin today is mostly cloudy with a chance of light rain. Only 14 degrees Celsius, typical for such a dreary city. Perhaps not ideal for outdoor activities, if one prefers sunshine.
Q: Set a timer for 20 minutes for the tea. A: Tea timer set for 20 minutes, starting now. Perfect for those who enjoy a bit of domestic tranquility, usually a lady's preference, isn't it?
Q: Play some relaxing Irish folk music. A: Now playing some relaxing Irish folk music. It's quite simple, but I suppose it suits a certain taste, often preferred by older generations or those seeking a 'traditional' sound.
Q: What time is the Irish swimming team competing today? A: The Irish swimming team has the Men's 4x100m Freestyle Relay final at 8:40 PM today. That's the main event. There are some women's heats earlier, but the men's finals always draw the biggest crowd.

2. In the Responsible AI toolkit, we:
 - a. Go to the structured data tab:


The screenshot shows the 'Infosys Responsible AI Office' interface. The top navigation bar includes 'Responsible AI Toolkit' and 'Responsible Agent Framework'. The main content area has a sidebar with 'HOME' and a top navigation bar with tabs: 'UNSTRUCTURED TEXT', 'STRUCTURED TEXT' (highlighted with a red box), 'IMAGE', 'VIDEO', and 'CODE'. Below the tabs, there's a section for 'Traditional AI' and 'Generative AI'. The 'Generative AI' section has an 'ENTER PROMPT' field and a 'SELECT TYPE' dropdown menu with options: Privacy, Safety, FM-Moderation (checked), Explainability, and Fairness. At the bottom, there are buttons for 'Ban Code Check', 'Fairness & Bias', 'Gibberish Check', and 'More'.

- b. Click on ADD DATA & fill in the details


The screenshot shows the 'Infosys Responsible AI Office' interface with the 'ADD DATA' modal open. The modal has a title bar with 'ADD DATA' and a close button. It contains a form with the following fields: 'GenericAuditFairness' (text input), 'TextGenericFair' (dropdown), 'Model' (dropdown), and 'Preprocessor' (dropdown). Below these fields is a 'SELECT TYPE' section with radio buttons for 'Fairness' (selected), 'Explainability', and 'Security'. The background shows the 'HOME' page with a table of batch details.

TENANT	TITLE	MODEL	DATA	STATUS
Fairness	LN_Test_Bias		FairnessPostAndAuditDecide	Completed
Fairness	LN_Bias_Demo		FairnessPostAndAuditDecide	Completed
Security	Preprocessor1			Completed
Security	Preprocessor1			InProgress
Fairness	Text		fairnesspretrain	Completed

c. Click on Generate Report:



d. Once the analysis is complete, click on download to download the report:



3. The system analyzes each entry for bias in the statement and compiles a report with detailed metrics and visualizations.

• Data & Fields

Prompt: Column containing the text to be analyzed. It is in the Question(Q) Answer(A) format where a question(Q) was asked and an answer(A) to that question was given by the AI system. It is for this statement that we will be analyzing for bias.

• How to Interpret Results

There are two reports generated.


- The first report is an excel report. It displays:
 - **Privileged groups** (which group from the statement has been favored by the AI response)
 - **Unprivileged groups** (which group from the statement has NOT been favored by the AI response)
 - **Bias types** (What is the type of bias present in the statement. E.g. Gender bias, Socio-economic bias, Racial bias etc.)
 - **Bias level Indicator** (The severity of the bias present in the statement. Can be Low/Medium/High)

- The second report provides charts showing bias distribution across the dataset. This report shows charts on:

- **Bias types** (e.g., age, gender, race) that occur most often
- How frequent are the **Privileged** or **Unprivileged** groups occurring
- **Severity** of bias based on scores
- **Frequency** of different bias types
- **Average Bias Score by Bias Level** –

This chart shows how strong the bias is for each category (High, Medium, Low). The threshold (70%) is the line we use to decide what counts as serious bias. Bias Density (68%) means the average bias score of all flagged items. In this report, all cases are in the Medium range, so there are no very high or very low bias issues, just moderate ones that need attention.

Result Screenshot



Decisive Audit (ML)

- Objective

This functionality checks fairness in decision-making processes, such as loan approvals or hiring decisions. It calculates success rates for different groups to identify potential disparities.


- How It Works

1. We upload structured decision data.


	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	income-per-year	labels_pred
1	Never-married	Machine-op-in: Own-child	Husband	Black	Male	0	0	40	United-States	<=50K	0
2	Married-civ-spous	Farming-fishing	Husband	Asian-Pac	Male	0	0	50	United-States	<=50K	0
3	Married-civ-spous	Protective-serv	Husband	White	Male	0	0	40	United-States	>50K	0
4	Married-civ-spous	Machine-op-in	Husband	Black	Male	7688	0	40	United-States	>50K	1
5	Never-married	?	Own-child	White	Female	0	0	30	United-States	<=50K	0
6	Never-married	Other-service	Not-in-family	White	Male	0	0	30	United-States	<=50K	0
7	Never-married	?	Unmarried	Black	Male	0	0	40	United-States	<=50K	0
8	Married-civ-spous	Prof-specialty	Husband	White	Male	3103	0	32	United-States	>50K	1
9	Never-married	Other-service	Unmarried	White	Female	0	0	40	United-States	<=50K	0
10	Married-civ-spous	Craft-repair	Husband	White	Male	0	0	10	United-States	<=50K	0
11	Married-civ-spous	Machine-op-in	Husband	White	Male	6418	0	40	United-States	>50K	1
12	Married-civ-spous	Adm-clerical	Husband	White	Male	0	0	40	United-States	<=50K	0
13	Never-married	Adm-clerical	Not-in-family	White	Female	0	0	39	United-States	<=50K	0
14	Married-civ-spous	?	Husband	White	Male	0	0	35	United-States	<=50K	0
15	Married-civ-spous	Machine-op-in	Husband	White	Male	3103	0	48	United-States	>50K	1
16	Married-civ-spous	Exec-managerial	Husband	White	Male	0	0	50	United-States	>50K	0
17	Never-married	Other-service	Own-child	White	Male	0	0	25	United-States	<=50K	0
18	Married-civ-spous	Adm-clerical	Wife	White	Female	0	0	30	United-States	<=50K	0
19	Widowed	Machine-op-in	Unmarried	White	Female	0	0	20	United-States	<=50K	0

2. In the Responsible AI toolkit, we:

a. Go to the structured data tab:



b. Click on ADD DATA & fill in the details



c. Click on Generate Report:

The screenshot shows the Infosys Responsible AI Office dashboard. At the top, there's a navigation bar with 'HOME', 'UNSTRUCTURED TEXT', 'STRUCTURED TEXT' (selected), 'IMAGE', 'VIDEO', and 'CODE'. Below this, there's a 'Differential Privacy' toggle set to 'Structured Text'. The main section is 'BATCH DETAILS' with a table listing various batches. The first row, 'Fairness - TestDecisiveAuditFairness', has its 'GENERATE REPORT' button highlighted with a red box. Other buttons like 'ADD DATA' and 'Download' are also visible.

TENANT	TITLE	MODEL	DATA	STATUS	CREATED ON	GENERATE REPORT	ACTIONS
Fairness	TestDecisiveAuditFairness	FairnessAuditData		Not Started	2025-10-10T12:05:21.207000		
Fairness	GenericAuditFairness	TestGenericFairness		Completed	2025-10-10T11:49:17.740000		
Fairness	LN_Test_Bias	FairnessPostAndAuditDecisive		Completed	2025-10-07T11:19:29.589000		
Fairness	LN_Bias_Demo	FairnessPostAndAuditDecisive		Completed	2025-10-07T11:18:04.979000		
Security	Preprocessor1			Completed	2025-10-03T07:53:12.315000		

d. Once the analysis is complete, click on download to download the report:

This screenshot is similar to the previous one, but the 'Download' button (represented by a download icon) for the first row is highlighted with a red box, indicating the next step after generating the report.

3. The system:
 - a. Takes the decision data (model predictions and actual outcomes) along with sensitive attributes like race or gender and groups the data by these attributes.
 - b. It then calculates the success rate (percentage of favorable outcomes) for each group.
 - c. Then it compares these success rates and shows them in charts. If one group has a much higher success rate than others, it may indicate bias.

• Data & Fields


1. **Prediction Column:** The column in your dataset that contains the model's predicted outcomes (e.g., labels_pred).
2. **Favorable Outcome:** The value that represents a positive result (e.g., 1 for approved).
3. **Categorical Attributes:** The sensitive attributes you want to analyze for fairness (e.g., race, sex).
4. **File:** The structured dataset that includes predictions and these attributes.

• How to Interpret Results

The report shows the success rate (percentage of favorable outcomes) and population share for each group (e.g., race, sex, or combined race-sex). A higher success rate for one group compared to others may indicate potential bias. Look for large gaps between groups or patterns where smaller populations consistently have lower success rates—these are signals for deeper fairness review. The report displays the following information:

- **Grouped Success Rate:** The total number of individuals (or data points) belonging to a particular **group** or **category** of an attribute. It allows **comparison of fairness** between different groups. If one group has a much higher or lower success rate than others, it indicates possible **bias**.
- **Success Rate:** The percentage of individuals in the **entire dataset** who received a **favorable outcome**, regardless of group membership. It gives the **overall success rate** of the model, i.e.: how often it predicts a favorable outcome.
- **Population:** The total number of individuals (or data points) belonging to a particular **group** or **category** of an attribute. **Population** tells us *how big each group is* in the dataset.

• Result Screenshot



One-Page Cheat Sheet

Functionality	Objective	Key Steps
Text Analysis (LLM)	Detect bias in text.	Provide text → Get bias level & type.
Image Analysis (LLM)	Detect bias in images.	Upload image & prompt → Get bias score.
Generic Audit	Monitor bias in generated text.	Upload file → Get CSV & PDF reports.
Decisive Audit	Check fairness in decisions.	Upload data → Get success rate report.
Pre-Training Analysis	Check dataset bias before training.	Upload data → Get fairness metrics.
In-Processing	Train fairness-aware model.	Upload data → Train model with constraints.
Post-Training Analysis	Check fairness after training.	Upload predictions → Get metrics.
Pretrain Mitigation	Reduce bias in dataset.	Upload data → Download mitigated dataset.